



C++ Modules: the good, the bad, the ugly

Uladzislau Chekhareu



- Technical Lead Developer at Solarwinds
- Love cats, C++ and riding a bike



- Modules - one of the most anticipated things in C++ 20
- But it's not discussed much comparing to some others!
- I'll try to explain, what is this beast and how to deal with it

The state of modules



- Paper [p1103r2.pdf](#) approved for merge into International Standard
- To some extent implemented by major compilers
- Examples from slides were tested using gcc

What do we have now



Organizing code through headers

- Old good `#include` directives
- Preprocessed, not part of the language
- Simple text replacement

What do we have now



Bad:

- Include guards
- Compilation times
- Weak physical encapsulation
- No isolation
- Sometimes include order matters
- ODR violations

What do we have now



Good:

- Easy
- Parallel compilation
- Good support by build systems
- Established practices to prevent smelling code

Modules goals



- Modular interfaces
- Physical encapsulation
- Isolation
- Compilation speed

New terms



`module`, `import` - new special identifiers

`export` - reused and repurposed keyword

Module unit - translation unit that contains module declaration

Module - collection of module units with the same name



Module units

Primary interface unit



```
// hello.cxx
export module hello;

export int foo() {
    return 42;
}
```

```
// main.cxx
import hello;

int main() {
    foo();
}
```

```
// valid module names
hello
hello.core
hello_util
hello101
```

```
// invalid module names
1111hello

.hello
```

Interface partition unit



```
// hello_world.cxx
export module hello:world;

export int bar() {
    return 50;
}
```



```
// hello.cxx
export module hello;
export import :world;

export int foo() {
    return 42;
}
```

Implementation unit



```
// hello_impl.cxx  
module hello;  
  
int foo() {  
    return 42;  
}
```



```
// hello.cxx  
export module hello;  
export import :world;  
  
export int foo();
```

Implementation partition unit



```
// hello_helper.cxx
module hello:helper;

int doStuff(int bar) {
    return bar + 1;
}

// hello_world.cxx ←
export module
hello:world;

import :helper;

export int bar() {
    return doStuff(50);
}

// hello_impl.cxx →
module hello;

import :helper;

int foo() {
    return doStuff(42);
}
```

Client code



```
// main.cxx
import hello;           // Ok
import hello:world;     // Error, can't explicitly import interface partition
import hello:helper;    // Error, can't import implementation partition

int main() {
    foo();               // Ok
    bar();               // Ok
    doStuff(1);          // Error
}
```



Exporting and importing



Export can only appear at namespace scope

```
export void foo() {} // Ok, global namespace

export template<typename T> // Ok
void foot(T input) { /* snip */ }

export enum class Errors { /* snip */ } // Ok, also allowed for old-style enums

namespace Bar {
    export int foo = 42; // Ok, namespace scope
    export struct FooBar { // Ok, also allowed for classes
        int Baz = 2;
    }
}
```



Export can only appear at namespace scope

```
class Gadget {  
    export Widget w;           // illegal  
};  
  
void bar() { export int i = 5; }    // bad  
  
void baz(export std::string input) {} // please don't  
  
template<export typename T>         // stop  
export class my_container{};  
  
enum class Errors {  
    export InvalidArgument = 1000    // error  
};
```

Convenient exporting



```
export namespace Goods { // exports all enclosing entities
    int const Count = 5;
    template<typename T>
    void foot(T input) { /* snip */ }
}
```

```
namespace Goods {
    void bar {} // NOT exported!
}
```

```
export {
    auto Value = 10.f;
    class Gadget {};
    /* more exported entities */
}
```

Can't export entities with internal linkage



```
namespace {  
    export int foo = 42;           // bad - anonymous namespace  
    export void bar() {}  
    export class MyClass() {};  
}  
  
export static int baz = 30;       // bad - static  
export static void foobar() {}
```

Can't export macros!



```
export #define FOO 42    // hell no!
```

```
#define BAR 42
```

```
export BAR;              // you shall not pass
```


Re-export of imported module



```
// gadget.cxx
export module gadget;

export struct Gadget
{
    int Detail;
};
```

```
// widget.cxx
export module widget;

// re-export
export import gadget;

export struct Widget {
    Gadget First;
};
```

```
// main.cxx
#include <iostream>
import widget;

using namespace std;

int main() {
    Gadget g{42};
    Widget w{g};
    cout << w.First.Detail << endl;
}
```

Modular code: all imports inside preamble



```
// foo.cxx
export module foo;
import bar;    // Ok
import baz;    // Ok

class baz {};  // Preamble ends, no more imports allowed

import foobar; // Error
```

Non-modular code: import everywhere



```
// main.cxx
#include <stdio.h>

import foo;           // Can use names from foo from this point

void foobar() {
    func_from_foo();  // Ok
    func_from_bar();  // Error
}

import bar;           // Can use names from bar from this point

void barfoo() {
    func_from_bar();  // Ok now
}
```


Can't import partitions belonging to other modules



```
// a_p1.cxx  
module a:p1;
```

```
/* snip */
```

```
// b_p1.cxx  
module b:p1;
```

```
/* snip */
```

```
// b.cxx  
export module b;  
import :p1;           // Ok  
import a:p1;          // Error  
import b;              // Error  
  
/* snip */
```

Cyclic import is not allowed



```
// a.cxx
```

```
export module a;
```

```
import b;
```

```
/* snip */
```

```
// b.cxx
```

```
export module b;
```

```
import c;
```

```
/* snip */
```

```
// c.cxx
```

```
export module c;
```

```
import a;           // not allowed
```

```
/* snip */
```

Partial backward compatibility



```
// foo.cxx
```

```
export module foo;
```

```
class import {};    // Ok, but please don't do it
```

```
import i1;          // Error, treated as keyword
```

```
::import i2;       // Ok, please don't
```



Modules in-depth

Name mangling



Before modules:

1. Internal linkage
2. External linkage
3. No linkage

Modules introduce:
4. Module linkage

Name mangling



```
// alinkage.cxx
```

```
export module alinkage;
```

```
export int foo = 42;    // external
```

```
static int baz = 50;    // internal
```

```
int bar = 10;           // module
```

```
export void doStuff() { // external
```

```
    int something = 5;    // no linkage
```

```
}
```

```
static void doInternalStuff() {} // internal
```

```
void doModuleStuff() {}           // module
```

Symbol table from alinkage.o:

Bind	Name
LOCAL	_ZW8alinkageEL3baz
LOCAL	_ZW8alinkageEL15doInternalStuffv
GLOBAL	foo
GLOBAL	bar
GLOBAL	_Z7doStuffv
GLOBAL	_ZW8alinkageE13doModuleStuffv

Name mangling



```
// alinkage_impl.cxx
// implementation unit of module alinkage
module alinkage;

void doBar() {
    bar = 110;           // Ok, module linkage
    baz = 150;           // Error, internal linkage
    doModuleStuff();    // Ok
    doInternalStuff();  // Error
}
```

Symbol table from alinkage.o:

Bind	Name
LOCAL	<code>_ZW8alinkageEL3baz</code>
LOCAL	<code>_ZW8alinkageEL15doInternalStuffv</code>
GLOBAL	<code>foo</code>
GLOBAL	<code>bar</code>
GLOBAL	<code>_Z7doStuffv</code>
GLOBAL	<code>_ZW8alinkageE13doModuleStuffv</code>

Name mangling



```
// blinkage.cxx
// separate module blinkage
export module blinkage;

export int foo = 50;    // Link time error
export int baz = 55;    // Ok
int bar = 15;          // Should be ok
void doModuleStuff() {} // Ok
```

Symbol table from alinkage.o:

Bind	Name
LOCAL	_ZW8alinkageEL3baz
LOCAL	_ZW8alinkageEL15doInternalStuffv
GLOBAL	foo
GLOBAL	bar
GLOBAL	_Z7doStuffv
GLOBAL	_ZW8alinkageE13doModuleStuffv

Name mangling implications



1. Modules are **not** name scoping mechanism!
2. Place exported names inside namespace
3. Use **static** and anonymous namespaces carefully



Headers

- Assuming m headers and n sources
- Source code size is $m + n$
- Header is compiled for each inclusion
- So compilation takes $m * n$ time

Modules

- Source code size is still $m + n$
- Module is compiled once
- Compilation takes $m + n$ time!
- How does it work?



BMI - binary module interface

- Produced by compiling module interface
- Read by importers
- Contains information required by importers to use module (for example, AST)
- Compiled once per build!



- BMI is produced by compiling module interface...
- And read by importers
- Dependency!
- Module interface **must** be compiled before importers
- Building modular code is not fully parallelizable
- Headers, on the other hand...



Paper concerning modular code compilation speed:

https://bfggroup.github.io/cpp_tooling_stats/modules/modules_perf_D1441R1.html

Author: Rene Rivera, committee member, participant of SG15

Twitter: <https://twitter.com/grafikrobot>



Method used:

- Compile 150 source files
- Test for various (up to 150) dependency DAG chain depths
- On each level include (or import) 3 components from previous levels
- Simple code - only *int* variables definitions

Compilation



Modular source

```
export module m148;

import m0;
import m15;
import m84;

namespace m148_ns
{
export int n = 0;
export int i1 = 1;
// ...
export int i300 = 300;
}
```

Header

```
#ifndef H_GUARD_h148
#define H_GUARD_h148
#include "h77.hpp"
#include "h78.hpp"
#include "h92.hpp"

namespace h148_ns
{
int n = 0;
int i1 = 1;
// ...
int i300 = 300;
}

#endif
```

Source

```
#include "h148.hpp"
```


Compilation



Jobs: 8

—▲— Non-Modular, GCC —●— Modular, GCC —▲— Non-Modular, Clang —●— Modular, Clang

GCC135: POWER9, altivec supported 2.2 (pvr 004e 1202) @ 2.166GHz (2 CPU, 32 cores, 128 threads)




Compilation



Jobs: 32

▲ Non-Modular, GCC ● Modular, GCC ▲ Non-Modular, Clang ● Modular, Clang

GCC135: POWER9, altivec supported 2.2 (pvr 004e 1202) @ 2.166GHz (2 CPU, 32 cores, 128 threads)




Compilation



Jobs: 128

—▲ Non-Modular, GCC —● Modular, GCC —▲ Non-Modular, Clang —● Modular, Clang

GCC135: POWER9, altivec supported 2.2 (pvr 004e 1202) @ 2.166GHz (2 CPU, 32 cores, 128 threads)




Compilation



Execution, Modular, Depth 20

GCC135: POWER9, altivec supported 2.2 (pic 004e 1202) @ 2.166GHz (2 CPU, 32 cores, 128 threads)




Compilation



Execution, Non-Modular, Depth 20


GCC135: POWER8, altivec supported 2.2 (pwr 004e 1202) @ 2.166GHz (2 CPU, 32 cores, 128 threads)





Observations:

- Unlike headers, modules build can not fully utilize system resources in highly parallel environments
- At least we can watch youtube while building cool modular code :)
- Moore's law
- But test is not perfect
- And compilers and build systems module support will mature





- `#include "legacy_header.h"` - inside global module fragment
- `import "legacy_header.h"` - inside module preamble
- Leaks macroses with all related problems
- Huge theme for another talk



Summary

Modules goals ?



- Modular interfaces
- Physical encapsulation
- Isolation (except exported names)
- Compilation speed (hopefully)

What do we do now?



- Wait and hope?
- Try it out yourself!
- Consider good practices for writing modular code (don't forget to share!)
- Prepare your legacy

How to try out



MSVC 2019

- `/experimental:module`
- `/std:c++latest`
- Modular code must be in ***.icc** files
- Only basic features are supported
- Somehow “supports” modular **stl** (*import std.core;*)

How to try out



GCC development branch 'c++-modules'

- Wiki link: <https://gcc.gnu.org/wiki/cxx-modules>
- -fmodules-ts
- -fmodule-header
- Mostly supports merging proposal
- Sometimes ICEs (specifically on some header units)

Some guidelines



Module naming:

- Module name should be unique
- Avoid names like **util** or **core**
- Prefix with company and/or product name



Exporting:

- All exports should be done within single top-level namespace
- Do not export multiple top-level namespaces
- Do not re-export everything, keep interfaces minimal

Conclusion



So are C++ Modules:

- Good?
- Bad?
- Ugly?



THANK YOU!



Q&A



- <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2018/p1103r2.pdf> - merging module proposal
- <https://gcc.gnu.org/wiki/cxx-modules> - GCC module implementation state
- https://bfgroup.github.io/cpp_tooling_stats/modules/modules_perf_D1441R1.html - compilation speed comparison
- <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2019/p1427r0.pdf> - module toolability concerns